

Image Processing

Lab 3

Kevin Gevers (s25595987)
Jeroen Overschie (s2995697)
Group 01

January 20, 2021

Note, that all used source code can be found attached next to this report's pdf file. Its structure should be self-explanatory; all requested functions are named accordingly and any extra functions are explained in the report. Note that (almost) every function has a corresponding test script, which is named just like its function, but with a suffix '_test'. Where possible, we followed the terminology from the book [1] for variable naming.

Exercise 1

First some basic theory on the subject. The goal is to work with Laplacian Pyramid decompositions, for which we apply both the decomposition process and the reconstruction process. The decomposition allows one to conveniently store images in multiple formats to use any suited version wherever desirable. By applying some filter, i.e. a Gaussian filter, we aid the subsampling process by making increasing the chance for picking pixels that are good representations of the original image; which the filter does by somewhat 'smoothing' the image. Moreover, by saving the differences between the various decomposition levels, we are able to reconstruct the original image by performing the same process in reverse, using the difference images to again obtain an image of the original dimensions.

Let us explain the notation used in the process. We are given some grey scale image, f . We now denote the Laplacian pyramid of f using:

$$f_1 = f$$
$$f_j = \text{REDUCE}(f_{j-1}) \text{ for } j = 2, \dots, J$$

Please do denote the subtle differences between the lowercase j and uppercase J . J denotes the number of levels of the pyramid composition, e.g. when $J = 3$, the pyramid will consist out of 3 images: the original, a decomposition of the original, and a decomposition of the decomposition. Furthermore, we define d_1, d_2, \dots, d_{J-1} as the *detail signals*, or *residuals*:

$$d_j = f_j - \text{EXPAND}(f_{j+1}) \text{ for } j = 1, 2, \dots, J-1$$

Which are more simply said the decompositions scaled back to the dimensions of their previous form, to then get the difference between the two images using matrix subtraction. Because we are diffing the images here, our vector d can only be of length $J-1$. Lastly, we define the REDUCE and EXPAND operators:

$$\text{REDUCE}(f) = \downarrow_2 (h_\sigma * f)$$
$$\text{EXPAND}(f) = h_\sigma * (\uparrow_2 (f)),$$

where h_σ denotes *Gaussian filtering*, which is a filter with a Gaussian response that can be used for performing a sort of image blur. The filter has the parameter σ , which is the standard deviation of the kernel. Furthermore, the operations denoted by the arrows, \downarrow_2 and \uparrow_2 denote **shrinking** and **zooming** an image by a factor 2, respectively. In the

spirit of maximum re-usability and efficiency, we reused some functions from Exercise 1 for this: the 'fundamental' functions `IPINTERPOLATE` (❖ Listing 2) and `IPSCALING_TRANSFORMATION` (❖ Listing 1) to do the scaling work, and `IPDOWNSAMPLE` (❖ Listing 3) and `IPZOOM` (❖ Listing 4) to apply them. These functions are thoroughly documented in the first report. In this lab, we use `IPDOWNSAMPLE` for our shrinking functionality (\downarrow_2) and `IPZOOM` for zooming functionality (\uparrow_2).

(a)

In this first exercise, we were asked to implement a function `IPPYR_DECOMP` that builds a Laplacian pyramid decomposition. For the parameters of this function, f is the input image, J the composition level and σ the Gaussian filter kernel parameter. For our implementation, see ❖ Listing 5.

In our implementation, we use matlab *Cell* types to store the pyramid and the detail signals, in variables f and d , respectively. Note that the Matlab cell type indeed makes it possible to store matrices of variable lengths, though we do need the function `MAT2CELL` and `CELL2MAT` for storing/retrieving from the cell array. Like announced, we were allowed to use the function `IMGAUSSFILT` to perform the Gaussian filtering process. We hand it a *double* image (`IM2DOUBLE`) and the parameter σ . We perform both the Reducing and Expanding operations and store the residuals in d .

Now finally, we can store the result in a matrix g . We can compute the height for this result matrix using:

$$P = M \times \left(1 + \frac{1}{2} + \left(\frac{1}{2}\right)^2 + \dots + \left(\frac{1}{2}\right)^{J-1}\right),$$

which we conveniently implemented using an element-wise squaring operation in Matlab, i.e. using `.^` (see `IPPYR_DECOMP` code line 35 in Listing 5). Note we add 1 to compensate for the Matlab indexing starting at 1 instead of 0. Next come some basic geometry operations to stack the result images vertically over the space of P and centering each image horizontally over the space of M . This done, we result with a result matrix g , which is also saved to `IPpyr_decomp-g_J=3, sigma=1.0.mat`.

(b)

Next, we visualize the results, using again a `*_test` script. See Listing 6 for the testing code. See Figure 1 for the original image and Figure 2 for its Laplacian pyramid decomposition using $J = 3$ and $\sigma = 1.0$. We can observe that the results look very much alike the example images given in the assignment instructions, despite having tuned the parameters slightly differently (sigma is higher in the assignment instructions, i.e. is 5.0 instead of our 1.0). As can be seen in the bottom most image, a decomposition of a 4x smaller size was made and still looks reasonably alike thanks to the filtering techniques applied at every step. Also the difference images look alike what was given in the assignment, indicating a correct decomposition.

Original 'plant' image




Figure 1: Original plant image for Exercise 1 without any filtering applied.




Figure 2: Laplacian pyramid decomposition of 1 using $J = 3$ and $\sigma = 1.0$. Residuals d are vertically stacked, with at the bottom the final decomposition, f_3 .

(c)

See `IPpyr_decomp_data-J=3, sigma=1.0.mat` for the result matrix data g .

Exercise 2

(a)

For exercise 2 we are going to reconstruct the decomposed image from exercise 1. The exercise states "Let f be a grey value image with Laplacian pyramid decomposition of J levels given by detail signals d_1, d_2, \dots, d_{J-1} and coarsest approximation f_J . As before you may assume that the image f is a square image of size $M \times M$, where M is a power of 2." To reconstruct the image we need to use Formula 1, which in turn needs Formula 2.

$$f_j = \text{EXPAND}(f_{j+1}) + d_j, \quad j = J-1, \dots, 1 \quad (1)$$

$$\text{EXPAND}(f) = h_\sigma * (\uparrow 2(f)) \quad (2)$$

Now we are tasked with writing a Matlab function `IPpyr_recon(g,J,sigma)` that implements the Laplacian pyramid reconstruction. Here g is the decomposed image that was the result of exercise 1, J is the number of decomposition levels, and σ is the standard deviation for the Gaussian filter. Our implementation can be seen in ♦ Listing 7.

(b)

Next we had to apply our function `IPpyr_recon(g,J,sigma)` to the result of exercise 1. We first load the .mat file that was saved in Exercise 1, extract the coarsest level image, f_J and then apply our reconstruction function `IPpyr_recon` to it. See ♦ Listing 8, for the final testing code.

Figure 3 shows the input image and the reconstructed image side by side. Here we used $J = 3$ and $\sigma = 1.0$ as stated in the assignment.



Figure 3: (Left) Input image. (Right) Reconstruction of the image after having decomposed it.

(c)

The next step is to calculate the mean absolute error between the reconstructed image and the input image. For this we need to use Formula 3 and we need to make sure both images are of the same type. We loaded the input file again as `uint8` and converted our reconstructed image to `uint8` as well.

$$\text{error}(f, g) = \frac{1}{M \times N} \sum_{i=1}^M \sum_{j=1}^N |f(i, j) - g(i, j)| \quad (3)$$

We also need to compute the difference image between the input and reconstructed images. This is done by taking the mean absolute value of each pixel value using both images. The result is presented in Figure 4.

(d)




Figure 4: (Left) Input image. (Middle) Reconstruction of the image after having decomposed it. (Right) Difference image between the input image and reconstructed image.

The last part of the assignment is to put the input image, the reconstructed image, and the difference image together in a single figure and save it. We did this as can be seen in ♦ Listing 8 and in Figure 4. The mean absolute error between the two images as implemented in (c) gives us the error 0, which means that the images are identical. This also shows in the completely black difference image in Figure 4. Using the Laplacian pyramid decomposition and reconstruction give the ability to compress the image without losing detail and later reconstruct the exact same image as the original.

Individual contributions

- **Exercise 1.** Contribution to program design, program implementation, answering questions posed and writing the report: *Jeroen* 100%.
- **Exercise 2.** Contribution to program design, program implementation, answering questions posed and writing the report: *Kevin* 100%.

References

- [1] Rafael C. Gonzalez and Richard E. Woods. *Digital image processing*. Prentice Hall, Upper Saddle River, N.J., 2008.

A Code

A.1 Exercise 1

A.1.1 Exercise 1 (a)

Listing 1: IPscaling_transformation.m: perform geometric transformations of the *scaling* kind, given an affine transformation matrix or a scaling constant.

```
1 function It = IPscaling_transformation(I, A, interpolation)
2 % IPscaling_transformation Computes an image scaling operation
3 % using an affine transformation matrix.
4 % Arguments:
5 %     I: Input image
6 %     A: Affine transformation matrix [3x3] of form
7 %         [cx 0 0; 0 cy 0; 0 0 1;] or a scalar value, then cx=cy.
8 %     interpolation: interpolation method. See IPinterpolate.m
9 if isnumeric(A)
10     A = [A 0 0; 0 A 0; 0 0 1;];
11 end
12 if ~exist('interpolation', 'var')
13     interpolation = 'none';
14 end
15 I = im2double(I);
16
17 % Image size
18 [M, N] = size(I); % height, width
19 D = [M, N, 1]; % dimensions
20
21 % Transformed dimensions
22 Dt = D * A;
23 Mt = round(Dt(1));
24 Nt = round(Dt(2));
25
26 % Map coordinates to new values
27 It = zeros(Mt, Nt);
28
29 %% Inverse mapping
30 % Perform _inverse mapping_ instead of forward mapping. Compute
31 % the inverse affine transformation matrix  $A^{-1}$ . See
32 % section 2.6 of (DIP, 42 – Gonzalez, Woods) book, page 102.
33 for y = 1:Mt
34     for x = 1:Nt
35         Pt = [x, y, 1];
36         P = Pt / A; % original coordinate (same as Pt * inv(A)).
37         offset = diag(0.5 * (1 - inv(A)))'; % inverse mapping centering offset
38         It(y, x) = IPinterpolate(I, P, offset, interpolation);
39     end
40 end
41
42 end
```

Listing 2: IPinterpolate.m: interpolate an unknown pixel intensity value using one of the supported methods.

```
1 function v = IPinterpolate(I, P, offset, interpolation)
```

```

2 % IPinterpolate Interpolate image pixel for image using given method.
3 % Arguments:
4 %     I: Input image
5 %     P: Coordinates of pixel with unknown intensity value; in original
6 %         (i.e. non-transformed) coordinate space. e.g. coordinates might
7 %         have decimals.
8 %     interpolation: one of
9 %         ('none' | 'nearest' | 'bilinear')
10 %         meaning no interpolation, nearest neighbor interpolation and
11 %         bilinear interpolation, respectively.
12 [M, N] = size(I);
13 switch interpolation
14     %% Nearest neighbor interpolation
15     % Finds the nearest pixel in the inverse mapping using 'ceil'.
16     case 'nearest'
17         % Make sure within bounds
18         Po = P + offset;
19         x = min(max(1, Po(1)), N);
20         y = min(max(1, Po(2)), M);
21         v = I(round(y), round(x));
22     case 'bilinear'
23         %% Bilinear interpolation
24         % See Bilinear interpolation Wikipedia page for used terminology and
25         % also https://bit.ly/3o3vcxD for parts of implementation.
26         Po = P + offset;
27         x = Po(1);
28         y = Po(2);
29         % Any values out of acceptable range
30         x(x < 1) = 1;
31         x(x > N - 0.001) = N - 0.001;
32         x1 = floor(x);
33         x2 = x1 + 1;
34         y(y < 1) = 1;
35         y(y > M - 0.001) = M - 0.001;
36         y1 = floor(y);
37         y2 = y1 + 1;
38         % Neighboring Pixels
39         Q = [I(y1,x1); I(y1,x2); I(y2,x1); I(y2,x2)];
40         % Pixels Weights
41         b = [(y2-y)*(x2-x); (y2-y)*(x-x1); (x2-x)*(y-y1); (y-y1)*(x-x1)];
42         v = dot(b, Q);
43     otherwise % no interpolation, i.e. 'none'
44         % if this pixel maps to an original pixel
45         if mod(P(1), 1) == 0 && mod(P(2), 1) == 0
46             v = I(P(2), P(1));
47         else % otherwise pad with zeros; no interpolation
48             v = 0;
49         end
50 end
51 end

```

Listing 3: IPdownsample.m: downsample images using IPSCALING_TRANSFORMATION.

```

1 function downsampledImage = IPdownsample(I, factor)
2 % IPdownsample Down-samples an image by a factor of 'factor', which

```



```

3 % must be a (positive) integer, i.e. factor >= 1.
4 % Arguments:
5 %   I: input image to downsample
6 %   factor: factor to shrink with. Must be an integer.
7 assert(isinteger(factor), 'factor must be an integer (was %d)', factor);
8 It = IPscaling_transformation(I, 1/double(factor), 'nearest');
9 downsampledImage = uint8(It * 2^8); % normalize back to 8-bit int image.
10 end

```

Listing 4: IPzoom.m: zoom images using IPSCALING_TRANSFORMATION.

```

1 function zoomedImage = IPzoom(I, factor)
2 % IPzoom Zooms an image by a factor of 'factor', which
3 % must be a (positive) integer, i.e. upsamplingFactor >= 1. Function zooms
4 % image by replicating pixels.
5 % Arguments:
6 %   I: input image to zoom
7 %   factor: factor to zoom with. Must be an integer.
8 assert(isinteger(factor), 'factor must be an integer (was %d)', factor);
9 It = IPscaling_transformation(I, double(factor), 'nearest');
10 zoomedImage = uint8(It * 2^8); % normalize back to 8-bit int image.
11 end

```

Listing 5: IPPYR_DECOMP function: Laplacian pyramid decomposition.

```

1 function g = IPpyr_decomp(f, J, sigma)
2 % IPpyr_decomp Laplacian pyramid decomposition
3 % Arguments:
4 %   f: input image
5 %   J: the number of decomposition levels
6 %   sigma: standard deviation of the Gaussian filter
7 f_1 = im2double(f);
8 [M, N] = size(f_1); % height, width
9 assert(M == N);
10
11 % Store pyramid and differences in cells
12 f = cell(1, J);
13 f(1) = mat2cell(f_1, M);
14 d = cell(1, J - 1);
15
16 % Build image pyramid
17 for j = 2:J
18     f_prev = cell2mat(f(j - 1));
19
20     % REDUCE
21     f_j = imgaussfilt(f_prev, sigma);
22     f_j = IPdownsample(f_j, uint8(2));
23
24     % EXPAND & difference
25     expanded = IPzoom(f_j, uint8(2));
26     expanded = imgaussfilt(expanded, sigma);
27     d_j = f_prev - expanded;
28
29     % Store in cells
30     f(j) = mat2cell(f_j, size(f_j, 1));

```

```

31     d(j - 1) = mat2cell(d_j, size(d_j, 1));
32 end
33
34 % Build output g
35 D = M * (1/2) .^ (0:J-1); % compute image dimensions D for all levels
36 P = sum(D) + 1; % +1 for Matlab indexing
37 x = M/2 - D/2 + 1; % compute 'g' x-coordinates
38 y = cumsum([0, D(1:J-1)]) + 1; % compute 'g' y-coordinates
39
40 % Insert detail signals 'd' and coarsest decomposition level 'J' into image
41 g = ones(P, M);
42 for j = 1:J
43     if (j == J); im = f(j); else; im = d(j); end
44     g(y(j):y(j)+D(j)-1, x(j):x(j)+D(j)-1) = cell2mat(im);
45 end
46 end

```

A.1.2 Exercise 1 (b)

Listing 6: IPpyr_decomp_test: Testing the Laplacian pyramid decomposition function, IPPYR_DECOMP.

```

1  clc; % clear the command window
2  close all; % close open figure windows
3
4  imname = 'plant';
5  inputfile = ['input.images/', imname, '.tif'];
6  f = imread(inputfile); % read input image
7
8  figure;
9  % Original image
10 colormap(gray(256));
11 imagesc(f);
12 axis equal;
13 axis tight;
14 title({'Original 'plant' image', ' '});
15 %%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
16 % Write current figure to file
17 all_file = ['output_plots/', imname, '_original', '.svg'];
18 saveas(gcf, all_file);
19 fprintf('\nComplete image has been saved in file %s\n', all_file);
20
21 % Pyramid decomposition
22 figure;
23 [M, N] = size(f); % height, width
24 J = 3;
25 sigma = 1.0;
26 g = IPpyr_decomp(f, J, sigma);
27 colormap(gray(256));
28 save('IPpyr_decomp_data-J=3,sigma=1.0.mat', 'g')
29 imagesc(g);
30 axis equal;
31 axis tight;
32 title({'pyramid decomposition', '(J=3, sigma=1.0)'});
33 %%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%

```

```

34 % Write current figure to file
35 all_file = ['output_plots/', imname, '_all', '_pyr-decomp', '.svg'];
36 set(gcf, 'PaperUnits', 'normalized')
37 set(gcf, 'PaperPosition', [0 0 0.75 1.00])
38 saveas(gcf, all_file);
39 fprintf('\nComplete image has been saved in file %s\n', all_file);

```

A.2 Exercise 2

A.2.1 Exercise 2 (a)

Listing 7: IPPYR_RECON function: Laplacian pyramid reconstruction.

```

1 function g2 = IPpyr_recon(g,J,sigma)
2 % IPpyr_recon Laplacian pyramid reconstruction
3 % Arguments:
4 %     g: result matrix from IPpyr_decomp
5 %     J: the number of decomposition levels
6 %     sigma: standard deviation of the Gaussian filter
7
8 [~, M] = size(g);
9
10 % calculate the coordinates needed for each level
11 D = M * (1/2) .^ (0:J-1); % compute image dimensions D for all levels
12 x = M/2 - D/2 + 1; % compute 'g' x-coordinates
13 y = cumsum([0, D(1:J-1)]) + 1; % compute 'g' y-coordinates
14
15 % extract coarsest level image
16 f_j = g(y(J):y(J)+D(J)-1, x(J):x(J)+D(J)-1);
17
18 for j = (J-1):-1:1
19     % EXPAND = upsample & Gaussian filter
20     f_j = IPzoom(f_j, uint8(2));
21     f_j = imgaussfilt(f_j, sigma);
22
23     % extract previous d
24     g_j = g(y(j):y(j)+D(j)-1, x(j):x(j)+D(j)-1);
25
26     % add previous d to expanded image
27     f_j = f_j + g_j;
28 end
29 g2 = f_j;
30 end

```

A.2.2 Exercise 2 (b-d)

Listing 8: IPPYR_RECON_TEST test script: executing the various task from exercise 2b-d.

```

1 clc; % clear the command window
2 close all; % close open figure windows
3
4 % Define J and sigma
5 J = 3;
6 sigma = 1.0;
7
8 % Load initial input image

```

```

9  imname = 'plant';
10 inputfile = ['input_images/', imname, '.tif'];
11 f = imread(inputfile); % read input image
12 f = im2double(f);
13 [M, N] = size(f);
14 assert(M==N);
15
16 % % Decompose the image
17 % g = IPpyr_decomp(f, J, sigma);
18
19 % Load decomposed image from file
20 mat = load('IPpyr_decomp_data-J=3,sigma=1.0.mat');
21 g = mat.g;
22
23 % Pyramid reconstruction
24 g2 = IPpyr_recon(g, J, sigma);
25
26 % Show results
27 figure;
28 subplot(121);
29 imshow(f)
30 subplot(122);
31 imshow(g2)
32
33 % Write current figure to file
34 saveas(gcf, 'output_plots/IPpyr_recon_test-partial.svg');
35
36 % Nicer output image but requires Matlab 2020a or above
37 % exportgraphics(gcf, 'output_plots/IPpyr_recon_test-partial.png');
38
39 % turn the next two lines off if you want to compare double images
40 g2 = im2uint8(g2);
41 f = imread(inputfile);
42
43 % compute absolute error between f and g2 + create difference image
44 diffImage = abs(f - g2);
45 error = sum(abs(f - g2), 'all') / (M * N);
46
47 % show results
48 figure;
49 subplot(131);
50 imshow(f)
51 subplot(132);
52 imshow(g2)
53 subplot(133);
54 imshow(diffImage)
55
56 % Write current figure to file
57 saveas(gcf, 'output_plots/IPpyr_recon_test.svg');
58
59 % Nicer output image but requires Matlab 2020a or above
60 % exportgraphics(gcf, 'output_plots/IPpyr_recon_test.png');

```